

Seminar Programmiersprachen im Vergleich - Zig

Tobias Schmitz

Contents

Introduction	3
Language	3
Development process	3
Single threaded optimization	3
Line by line matching	3
Whole text matching for performance	3
Line by line matching with fixed size buffer	3
Whole text matching for performance with fixed size buffer	3
Parallelization	4
Conclusion	4

Introduction

Language

- build system
 - same language for build-system
 - build.zig is generated
 - good C interoperability
- manual memory management
 - allocators have to be manually passed
 - deallocate by deferring `deinit` procedures
 - memory leaks are automatically detected in debug mode
- minimal standard library
 - no unicode string support, only functions for operating on slices e.g. `std.mem`
 - utf-8 string libraries are 3rd party
- exhaustive switch statements
 - useful for enums
- generics are just compiletime functions operating on types
- inferred struct literal type `.{}`
- if enum type is known the type can be omitted, `.variant` is sufficient
- errors as values
 - error union explicit or implicit
 - try: return on error
 - catch: handle error
- slices
 - pointer and length by default: `[]u8`
 - sentinel terminated, for example null terminated: `[*0]u8`
 - exclusive ranges for slicing: `my_slice[0..2]`
- somewhat immature ecosystem
 - missing regex library
 - async not available in 0.11 self-hosted compiler
- unclear crash messages, even in debug mode
 - just memory addresses
 - this is a bug when linking `libc`

Development process

Single threaded optimization

Line by line matching

Whole text matching for performance

- to avoid allocating iterator for every line
- then avoid using iterator anyway

Line by line matching with fixed size buffer

- no multiline matches
- easier to implement
- lines need to be iterated anyway for line numbers

Whole text matching for performance with fixed size buffer

- the rust regex find functions has extremely high startup overhead on unicode word-character patterns
- avoided to some extent by searching the whole text buffer instead of just line by line slices

Parallelization

- Two worker thread pools
 - One for walking the file tree
 - One for searching files
- Walking the filetree uses a shared stack for depth first searching
- A shared queue is used to pass to be searched files
- A shared sink synchronizes writing to stdout

Conclusion